

Upwatch User Guide

Ron Arts

Upwatch User Guide
by Ron Arts

Copyright © 2002-2004 by UpWatch BV, The Netherlands

Table of Contents

Preface.....	i
1. About UpWatch	1
1.1. History.....	1
1.2. Features.....	1
2. Installation	2
2.1. Getting upwatch.....	2
2.2. Requirements.....	2
2.2.1. Run-time requirements.....	2
2.2.2. Build requirements.....	2
2.3. Compiling upwatch.....	2
3. Configuration	4
3.1. Configuration Files.....	4
3.2. Monitoring specific log files.....	5
4. Utilities	7
4.1. mbmon.....	7
4.2. chklog.....	7
4.3. uwq.....	7
4.4. uwregexp.....	7
4.5. uwsaidar.....	7
Index.....	8

Preface

This book describes the client side installation and configuration of UpWatch.

Chapter 1. About UpWatch

1.1. History

UpWatch is born from the loins of Netland Internet Services BV, Amsterdam, The Netherlands. We are a hosting company which started in 1993 (when even Bill Gates knew nothing of the internet). We started doing managed hosting in 1995, and that's when we found out about monitoring. The hard way.

It became clear that customers can easily bring down their own server on impossible moments, and that it doesn't look very good if you both find out about that the Monday after. So we started doing SLAs and limit customer rights on their own server.

Initially we used Big Brother (bb4.com) for monitoring. This is an outstanding and useful package, and we have been using it for many years. But it has a few downsides. One is scalability. It does not scale well to hundreds of hosts. Also it has a geek-like look, we felt we couldn't give the URL to our customers. Third problem was integration with our backoffice.

At the same time yours truly was thinking about setting up a commercial service for monitoring servers remotely. All this culminated into UpWatch. So let's get straight to the ..

1.2. Features

This is the full list of relevant upwatch features:

- OS support: clients available for Linux, Windows, FreeBSD, Solaris
- GUI is multi-language enabled (uses gettext)
- GUI has mobile client support
- GUI is brandable, you can give it your own look & feel
- Generates realtime graphs from the database
- Notifications by email or SMS.
- Clients for: HTTP GET, IMAP, MSSQL, MySQL, PING, POP3, PostgreSQL, SMTP, SNMP GET, TCP connect (any port)
- Local client detects: CPU load, loadavg, swap use, I/O use, memory use, and where supported hardware info like CPU temperature, fan speed and Power voltages. Also you can set it up to scan any logfile using regular expressions you supply.
- Scalable: designed for monitoring tens of thousands of hosts
- Multi-tenanting: multiple companies can run monitoring services for network of multiple client-companies using the same backend+probe serverpark
- Extensive and complete documentation, partly generated from source
- Secure: run as ordinary user, developed with security in mind
- Fully opensource: GUI built on Apache/PHP, Backend on C/Perl, Database is MySQL. Uses GNU configure.
- SuSE, RedHat and Fedora RPM's generated from sourcetree for easy installation

Chapter 2. Installation

2.1. Getting upwatch

The UpWatch clients can be download from www.upwatch.com. They are available in the following formats:

- as a tar.gz file, including sources for every supported platform
- SuSE, RedHat or Fedora .rpm files
- A windows installer

If you want, you can inspect the code for security issues.

2.2. Requirements

2.2.1. Run-time requirements

First ensure that the time/date on all hosts is correctly set.

Here's a list of everything we expect (I'll also list the version we use ourselves):

- glib2 >= 2.0.4
- xml2 (any version will do)
- libpcrc 3.9.10
- libncurses 5.2
- libreadline 4.3

Delivered with upwatch are libstatgrab (0.10) (<http://www.i-scream.org/libstatgrab/>), to retrieve critical operating system values, xmbmon 2.03 (<http://www.nt.phys.kyushu-u.ac.jp/shimizu/download/download.html>) for hardware statistics, and the State Threads Library (1,4) (<http://state-threads.sourceforge.net/>) for fast and efficient multithreading.

2.2.2. Build requirements

You probably don't want to build upwatch yourself. Most likely you'll grab the RPM packages and issue `rpm -Uvh upwatch*.rpm`. Then skip configuration

But on the other hand: you can build the software yourself. Apart from the normal GNU compilation tools, and the development versions of the above mentioned packages, you'll need the following on your system to build upwatch:

- autogen 5.3.6 (autogen.sourceforge.net)
- RPM tools, if you want to build RPM's

If you run RedHat, Debian or SuSE, don't forget to install the *-devel packages if there are any.

2.3. Compiling upwatch

Just in case you really want to (or need to) compile upwatch yourself, it's pretty easy:

```
$ tar xzvf upwatch-x.x.tar.gz
$ cd upwatch-x.x
$ ./configure
$ make
$ make install
```

Nothing to it... In case of problems, you're probably missing some library or header files, or they are in unexpected places. Look at the last parts of config.log.

Chapter 3. Configuration

3.1. Configuration Files

The directory structure for configuration files is as follows. At the top level is `upwatch.conf` (usually residing in `/etc`). This file is read by all programs, and contains global parameters, and parameters you want to make globally known. At the same level is the directory `upwatch.d`. This contains configuration files for every program. The `upwatch.conf` file looks like this:

```
# Upwatch configuration file
# contains defaults for all modules
# these can be overridden in /etc/upwatch.d/<module>.conf
#

debug 2
syslog no
stderr no

logfile /var/log/upwatch/messages

spooldir /var/spool/upwatch
```

All values in this file can be overridden or augmented in program-specific files. `debug` speaks for itself. Never set it higher than two. Zero suppresses all debugging except the most critical ones, 'debug 1' will output only warnings, 'debug 2' will send progress information to the log. `syslog yes` will enable logging to the system log. `stderr` is really not very useful because the commandline parameter `-e` accomplishes the same for every program. `logfile` value denotes where logging will take place. This file should be writable by the user `upwatch`. The same holds for the `spooldir` base directory. It should contain subdirectories of the maildir format (meaning: each having a `andtmp` subdirectory).

On the average client two processes will be running continuously: `uw_sysstat` and `uw_send`. The first program collects info on your system, and writes an XML file, the second sends it to the central database. `upwatch.d` will contain the files `uw_sysstat.conf` and `uw_send.conf`. Let's first look at `uw_send`. It looks like this:

```
# where to send to
host cms-db.of.ce.netland.nl
port 1985
uwuser 20010631
uwpasswd SaSNF8bu
debug 2
threads 1
# where to read from
input uw_send
```

Host and port refer to the central database location. You need user and password to log into that. The `debug` setting determines the amount of logging that the program does, and `threads` how many files to send at the same time, and finally, `input` tells `uw_send` where to find its XML files to send out. Pretty straightforward.

The `uw_sysstat.conf` file is comparable:

```
realm netland
serverid 381
output uw_send
hwstats on
errlog syslog /var/log/messages
errlog maillog /var/log/maillog
```

The `realm` is basically a short name for your company. It tells the central processor in which database to store data from this server. `serverid` is the numerical id of this server in that particular database. Without it it would not know where to store the XML result. `hwstats` determines (where supported) if `uw_sysstat` will try to talk to the motherboard to get hardware values like temperature and fanspeed (module must be present to test if this is supported) and the `errlog` parameters tell `uw_sysstat` which files to monitor and their format.

3.2. Monitoring specific log files

`uw_sysstat` is special in that it allows you to scan every (line-oriented) log file you want. It uses regular expressions to set a yellow or red state. It works as follows:

On startup it reads `/etc/upwatch.d/uw_sysstat.conf`, and searches for log file statements. Say it encounters the statement:

```
log file errlog /var/log/messages
```

what it does is it reads all files in the directory `/etc/upwatch.d/uw_sysstat.d/errlog` (except `rmacros.txt` and `macros.txt`). These files should contain regular expressions prefixed by one of the keywords `green`, `yellow`, or `red`. Next `uw_sysstat` starts scanning `/var/log/messages`. It reads a line from the log file and the following happens:

- Check against the red list. If match found, flag red condition, and send the offending line to the upwatch server
- Check line against the yellow list. If matches, flag yellow and send to server
- Check against green list. If it matches, ignore this line and go the next line in the log file. If the current line does not match any of the green list, flag yellow, and send line to server

The regular expressions may (for readability) contain macros, they should be entered in `/etc/upwatch.d/uw_sysstat.d/syslog/macros.txt`.

You can easily add a directory of your own, containing regular expressions for your own log files. In fact upwatch includes a handy utility `chklog` to help you create regular expression lists. Here is an example how to do it.

Suppose you plan to scan the log file for the imaginary 'timtim' navigational system. It resides in `/var/log/timtim.log`.

- First create the directory:


```
# cd /etc/upwatch.d/uw_sysstat.d
# mkdir timtim
# cp syslog/rmacros.txt timtim
```

```
# cp syslog/general timtim
# chown -R root:upwatch timtim
# chmod 770 timtim
# chmod 660 timtim/*
```

- Next look at macros.txt and tailor it to accommodate specifics for the timtim logfile. It might for example contain entries for zipcodes, or latitudes/longitudes for which you would like to create macros.

- Next step: extract regular expressions from an example logfile:

```
# chklog -t timtim -r /var/log/timtim.log | sort -u > /tmp/timtim
```

- edit this file. Maybe it will contain dupes, you should try to keep the number of regular expressions low. If you are satisfied you can try copying it to /etc/upwatch.d/uw_sysstat/timtim using any filename, and use chklog to test it:

```
# chklog -t timtim -m /var/log/timtim.log
```

Now you should only see the lines you want to be reported by uw_sysstat. Repeat steps until you are satisfied

- Finally tell uw_sysstat that you want it to start scanning by adding timtim /var/log/timtim.log to its configuration file. That's it.

Chapter 4. Utilities

4.1. mbmon

On x86 architectures and some operating systems hardware readouts may be obtained. We use xmbmon (<http://www.nt.phys.kyushu-u.ac.jp/shimizu/download/download.html>) to get these values, so wherever that's supported, we can do it. But before switching on hwstats in uw_sysstat, use mbmon to test if your setup is supported. By the way, if mbmon is not on your system, it's definitely not supported. anyway, also look if the values given by mbmon are meaningful.

Each program has a manual page that documents options. Every long commandline option can also be entered in a configuration file.

4.2. chklog

This utility makes it simpler to create your own set of regular expressions for a particular log file you want to be monitored by uw_sysstat. It can both scan a test log file and output regular expressions, as scan a log file and outputs lines that are suspicious, and to which uw_sysstat should respond by adding a yellow or red condition. See Monitoring specific log files for the procedure.

4.3. uwq

You can show all queues on the current system with uwq. We regularly use uwq to monitor queue status.

4.4. uwregex

While creating a regular expression for an application specific log file, sometimes you get problems creating a regular expressions that it's really well. This is where uwregex comes in handy. You feed it the line to match, and the log file type, and it gives you a prompt where you can try various regular expressions. Use up and down arrows to cycle through previous expressions.

4.5. uwsaidar

this is a general handy tool, a bit like top, it gives you an ongoing system status screen with CPU, memory, and I/O usage and various other parameters. It actually comes straight out of the libstatgrab (<http://www.i-scream.org/libstatgrab/>) library, that accompanies the Upwatch client.

Index